


# CS 4390

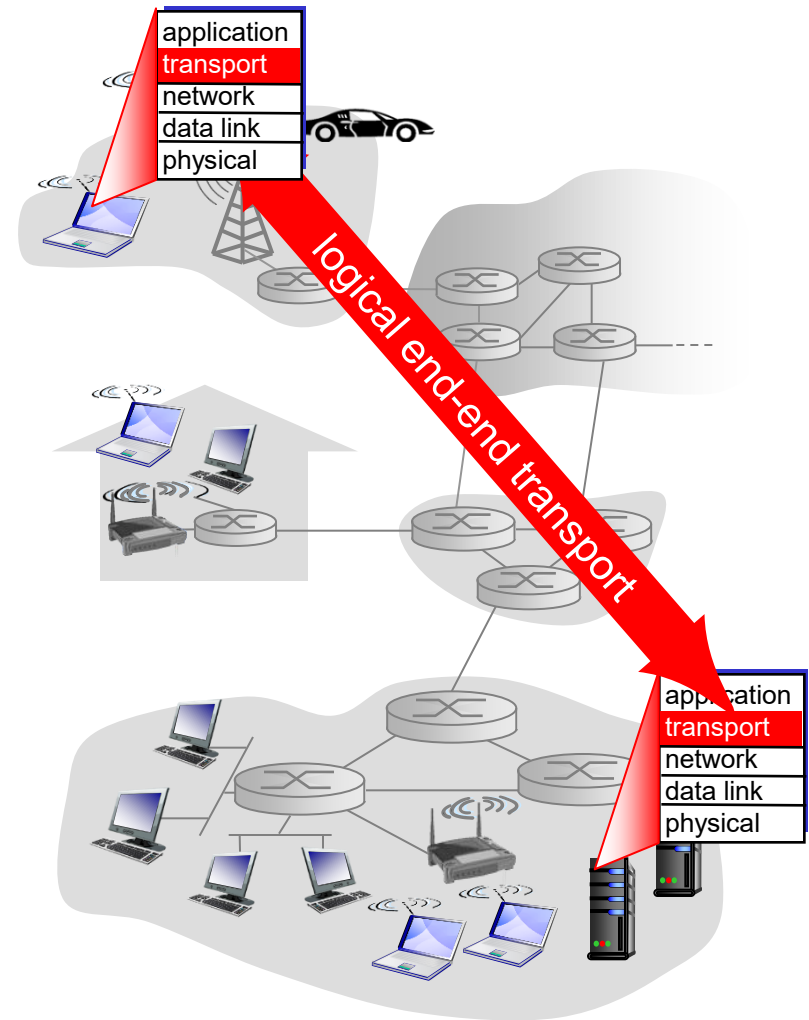
## Computer Networks



### *Transport Layer – Overview and UDP*

# Transport Services and Protocols

- ❖ provide *logical communication* between app processes running on different hosts
  - Apps require different *services*
- ❖ transport protocols run in end systems
  - *send side*: breaks app messages into *segments*, passes to network layer
  - *receive side*: reassembles segments into messages, passes to app layer
- ❖ more than one transport protocol available to apps
  - Internet: TCP and UDP



# Challenges for Transport Layer

- Develop ***algorithms*** and ***techniques*** for transport protocols that address ***transport requirements*** imposed by application programs for
  - Datagram service – connection-less
  - Byte stream service – connection-oriented
- Take into account ***limitations*** of underlying layers (network and below)
  - Network layer services and protocol limitations
  - Underlying physical data transmission network limitations, e.g. higher bit-error rate in wireless channels

# Transport Service Requirements

- Supports multiple application processes on a host
- Supports arbitrarily large messages
- Guarantees message delivery
- Guarantees no errors in delivered messages
- Delivers at most one copy of each message
- Delivers messages in the same order they were sent
- Ensures that messages not overwhelm the receiver
- Delivers messages with minimum delay
- ...

# Limitations of Underlying Layers

- Bit-errors in messages
  - Vary depending on physical channels, can be high in wireless
- Messages are lost
  - e.g. discarded by routers due to buffer overflow
- Messages are delivered out-of-order
  - Packet switching: packets may traverse different paths from sender to receiver
- Duplicated messages
  - Duplicated copies of the same messages may be received (due to retransmission)
- Limit messages size
  - Ethernet frames: 1.5 KB, IP: 64 KB
- Delayed delivery due to congestion
- ...

# Transport Layer Topics

## our goals:

- understand *principles* behind transport layer services:
  - multiplexing, de-multiplexing
  - reliable data transfer
  - flow control
  - congestion control
- learn about Internet transport layer protocols:
  - UDP: connectionless transport
  - TCP: connection-oriented reliable transport
  - TCP congestion control

# Transport v.s. Network Layer

## ❖ *network layer:*

logical  
communication  
between *hosts*

## ❖ *transport layer:*

logical  
communication  
between *processes*

- relies on, enhances,  
network layer  
services

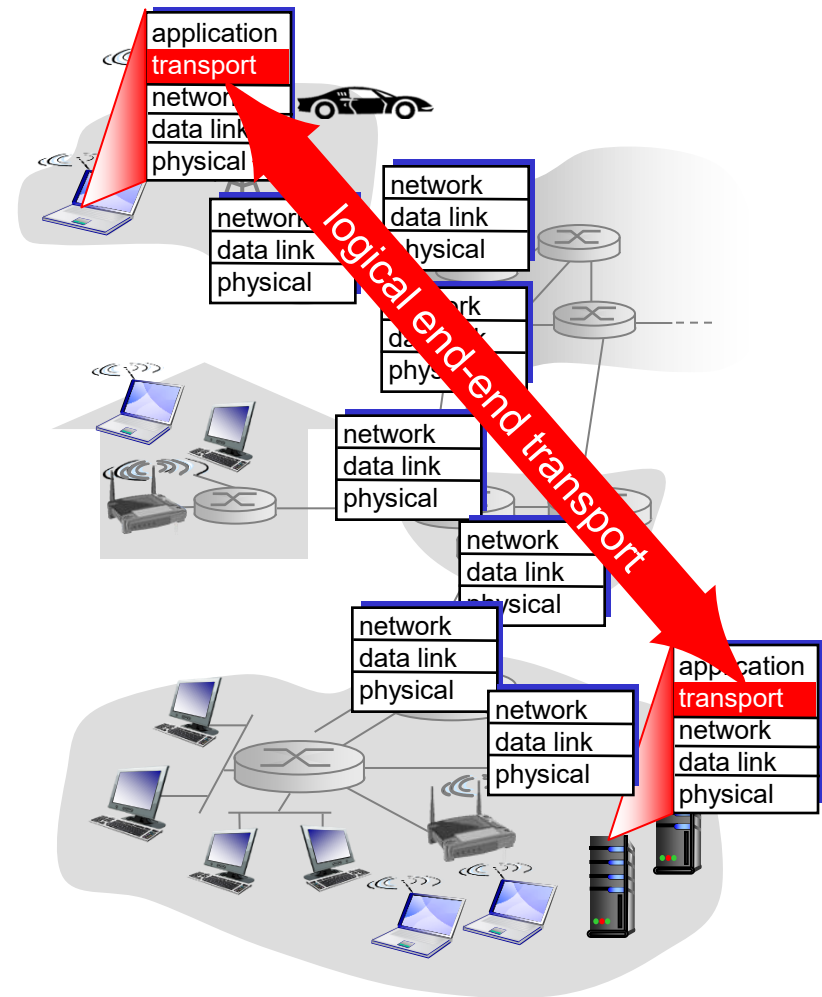
## *household analogy:*

*12 kids in Ann's house  
sending letters to 12 kids in  
Bill's house:*

- hosts = houses
- processes = kids
- app messages = letters in envelopes
- transport protocol = Ann and Bill who de-mux to in-house siblings
- network-layer protocol = postal service

# Internet Transport-layer Protocols

- reliable, in-order delivery: **TCP**
  - congestion control
  - flow control
  - connection setup
- unreliable, unordered delivery: **UDP**
  - no-frills extension of “*best-effort*” IP
- services not available:
  - delay guarantees
  - bandwidth guarantees





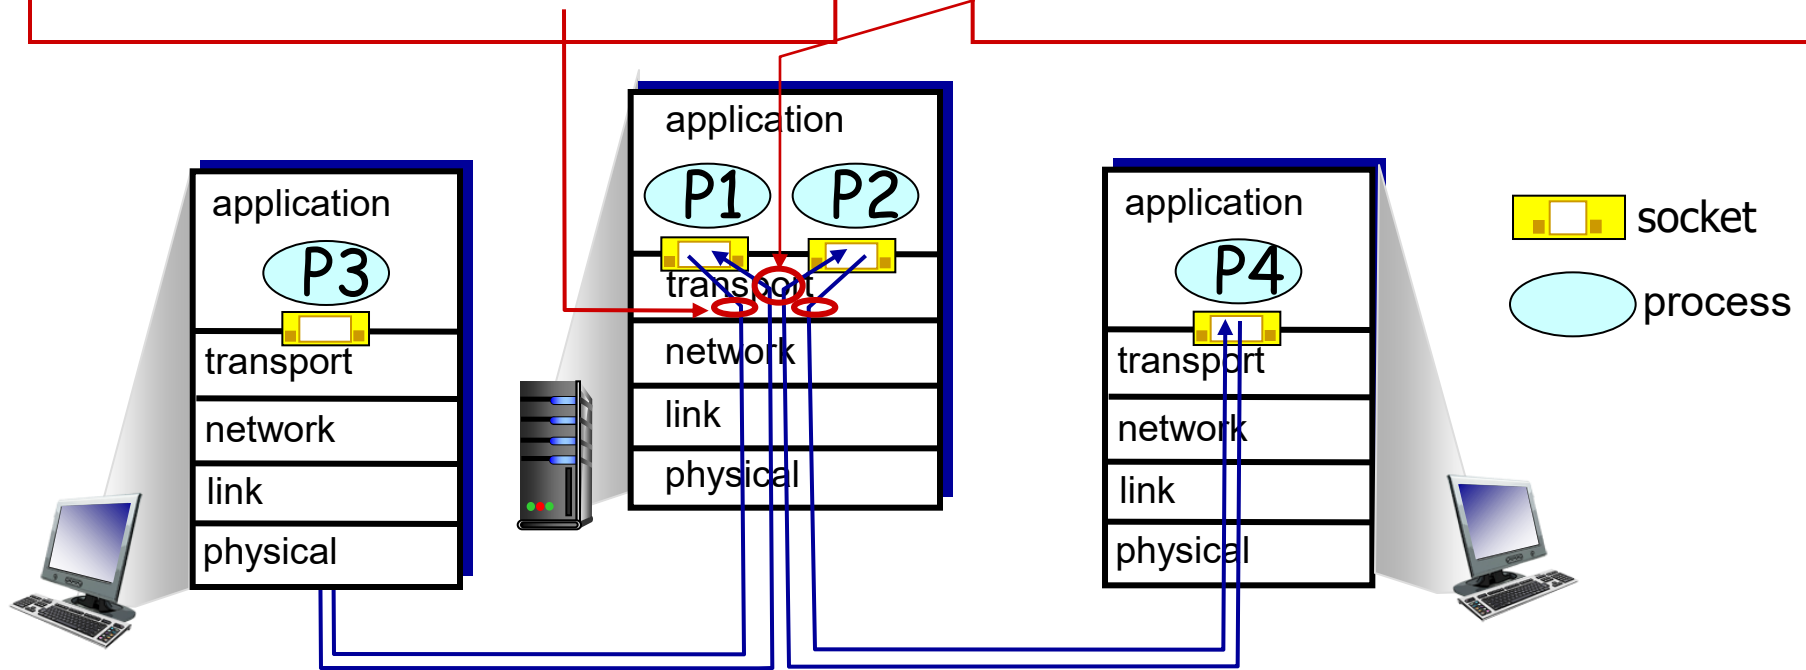
# Multiplexing & Demultiplexing

## *multiplexing at sender:*

handle data from multiple sockets, add transport header (later used for demultiplexing)

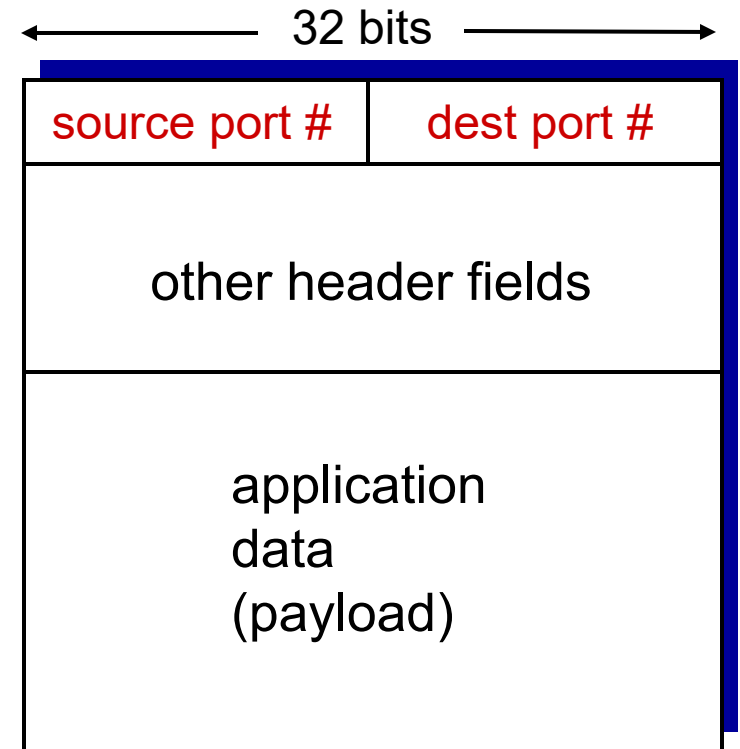
## *De-multiplexing at receiver:*

use header info to deliver received segments to correct socket



# How De-multiplexing Works?

- ❖ host receives IP datagrams
  - each datagram has source IP address, destination IP address
  - each datagram carries one transport-layer segment
  - each segment has source, *destination port number*
- ❖ host uses *IP addresses & port numbers* to direct segment to appropriate socket



TCP/UDP segment format

# Connectionless De-multiplexing

- *recall*: created socket has host-local port #:

```
DatagramSocket mySocket1  
= new DatagramSocket(12534) ;
```

- ❖ *recall*: when creating datagram to send into UDP socket, must specify
  - destination IP address
  - destination port #

- 
- ❖ when host receives UDP segment:

- checks destination port # in segment
- directs UDP segment to socket with that port #



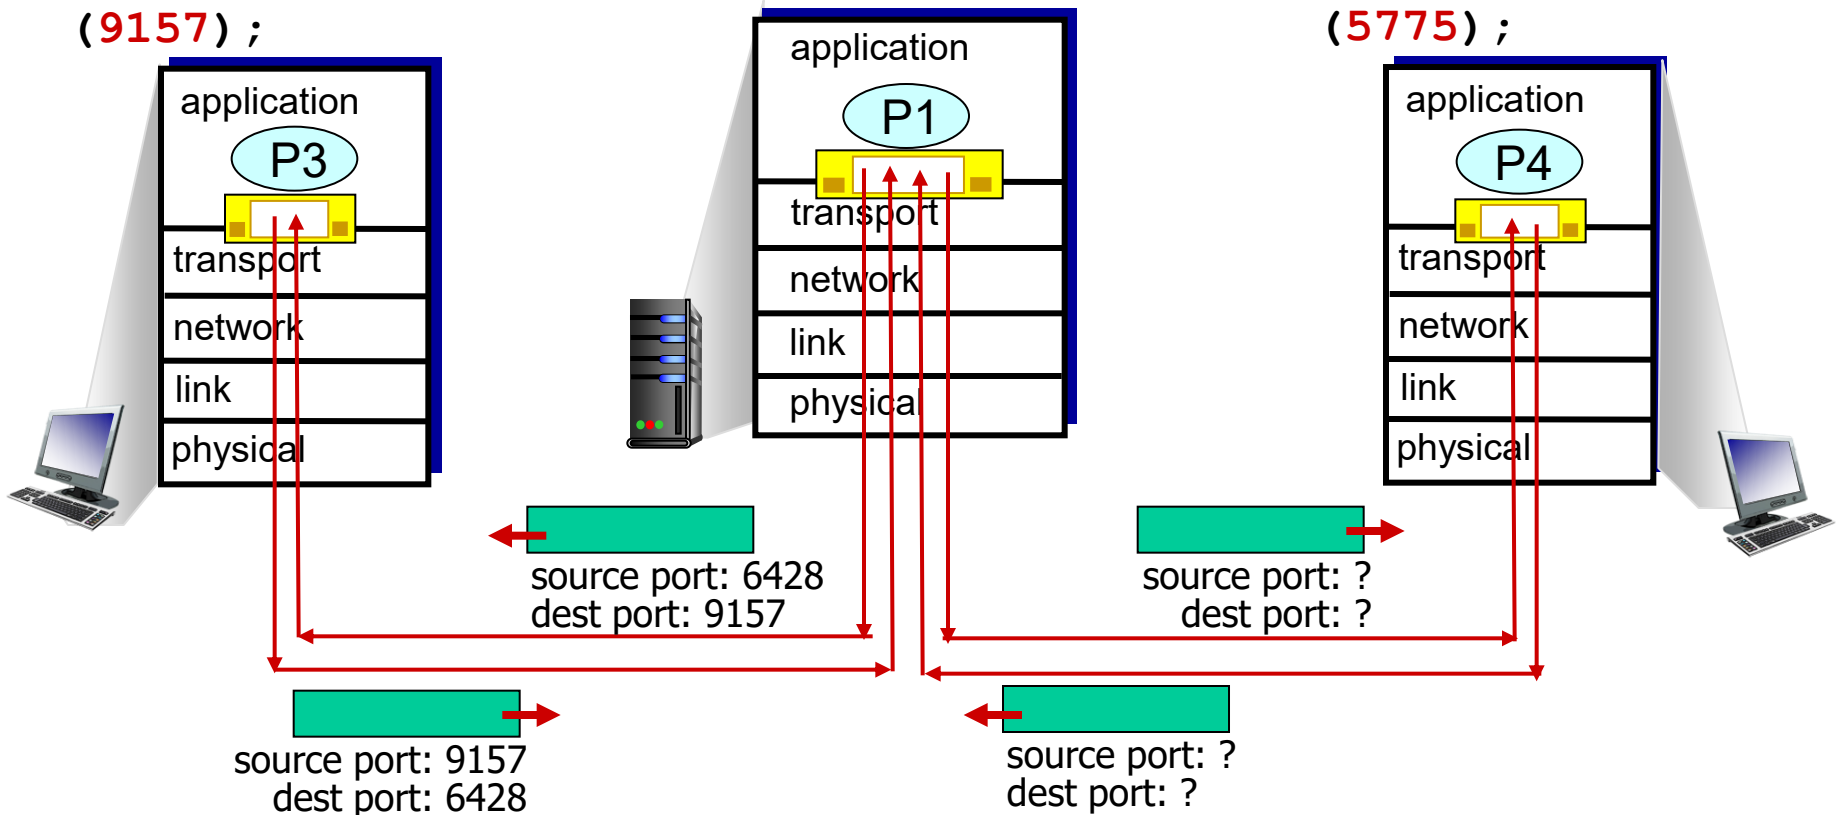
IP datagrams with *same dest. port #*, but different source IP addresses and/or source port numbers will be directed to *same socket* at destination

# Connectionless De-mux – Example

```
DatagramSocket  
mySocket2 = new  
DatagramSocket  
(9157) ;
```

```
DatagramSocket  
serverSocket = new  
DatagramSocket  
(6428) ;
```

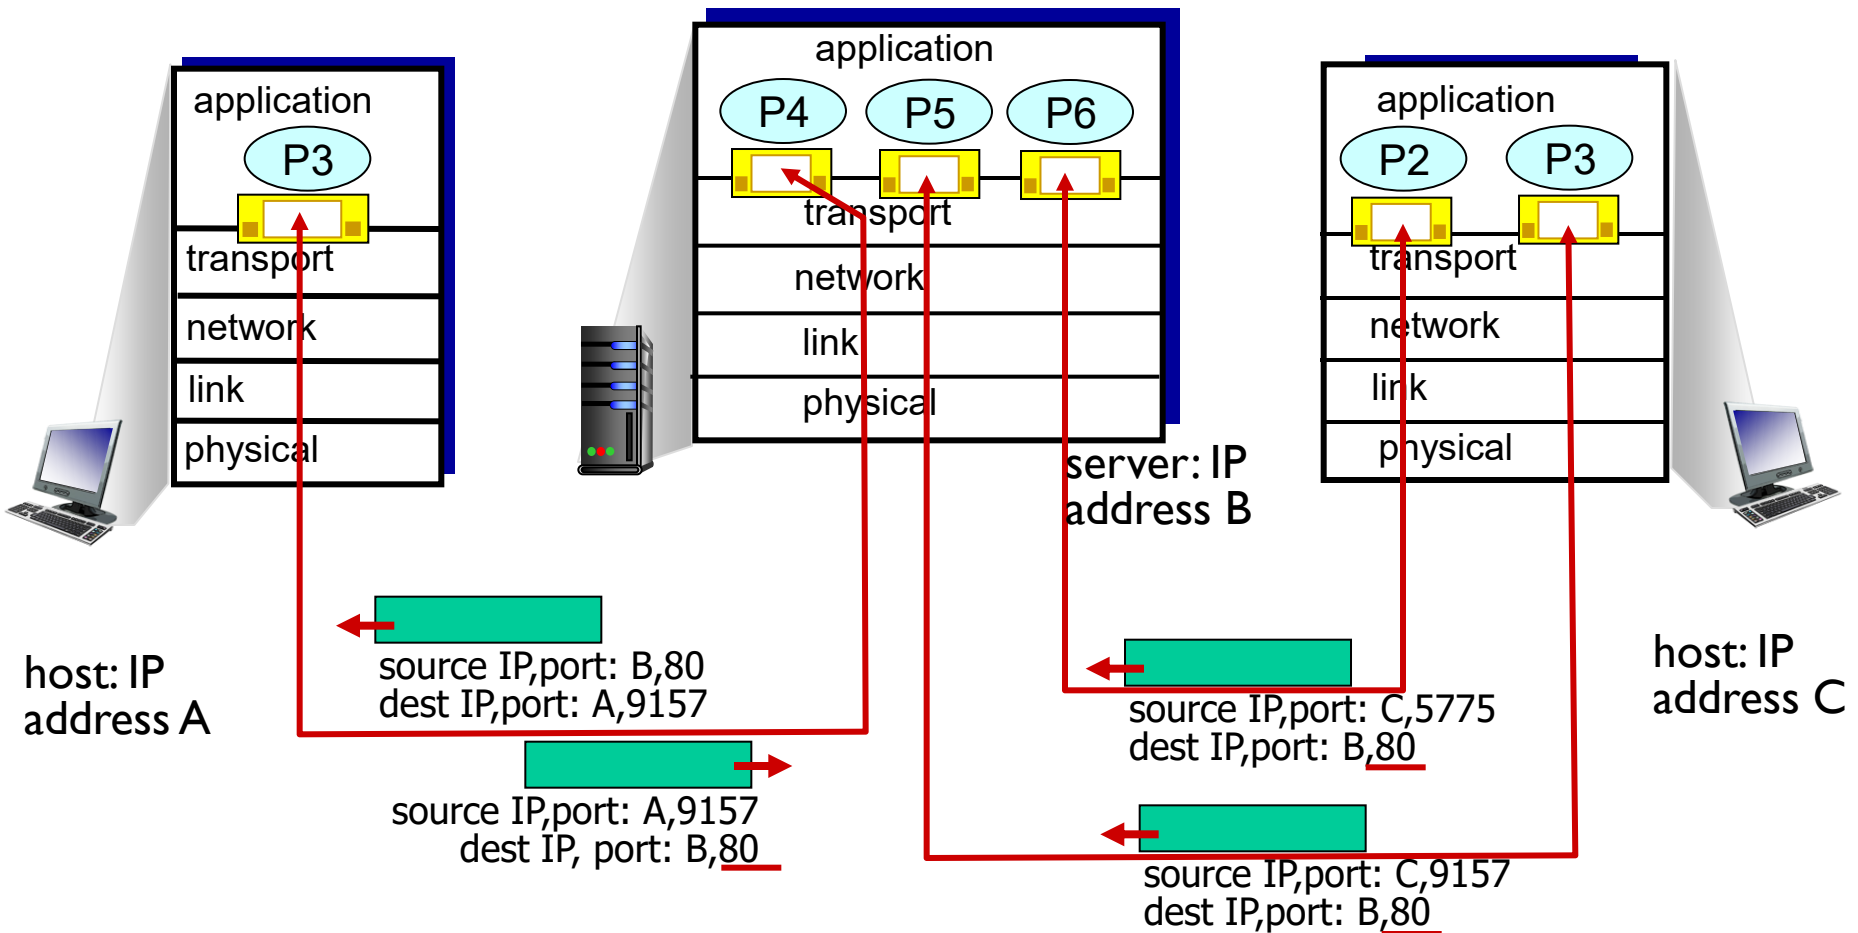
```
DatagramSocket  
mySocket1 = new  
DatagramSocket  
(5775) ;
```



# Connection-oriented De-mux

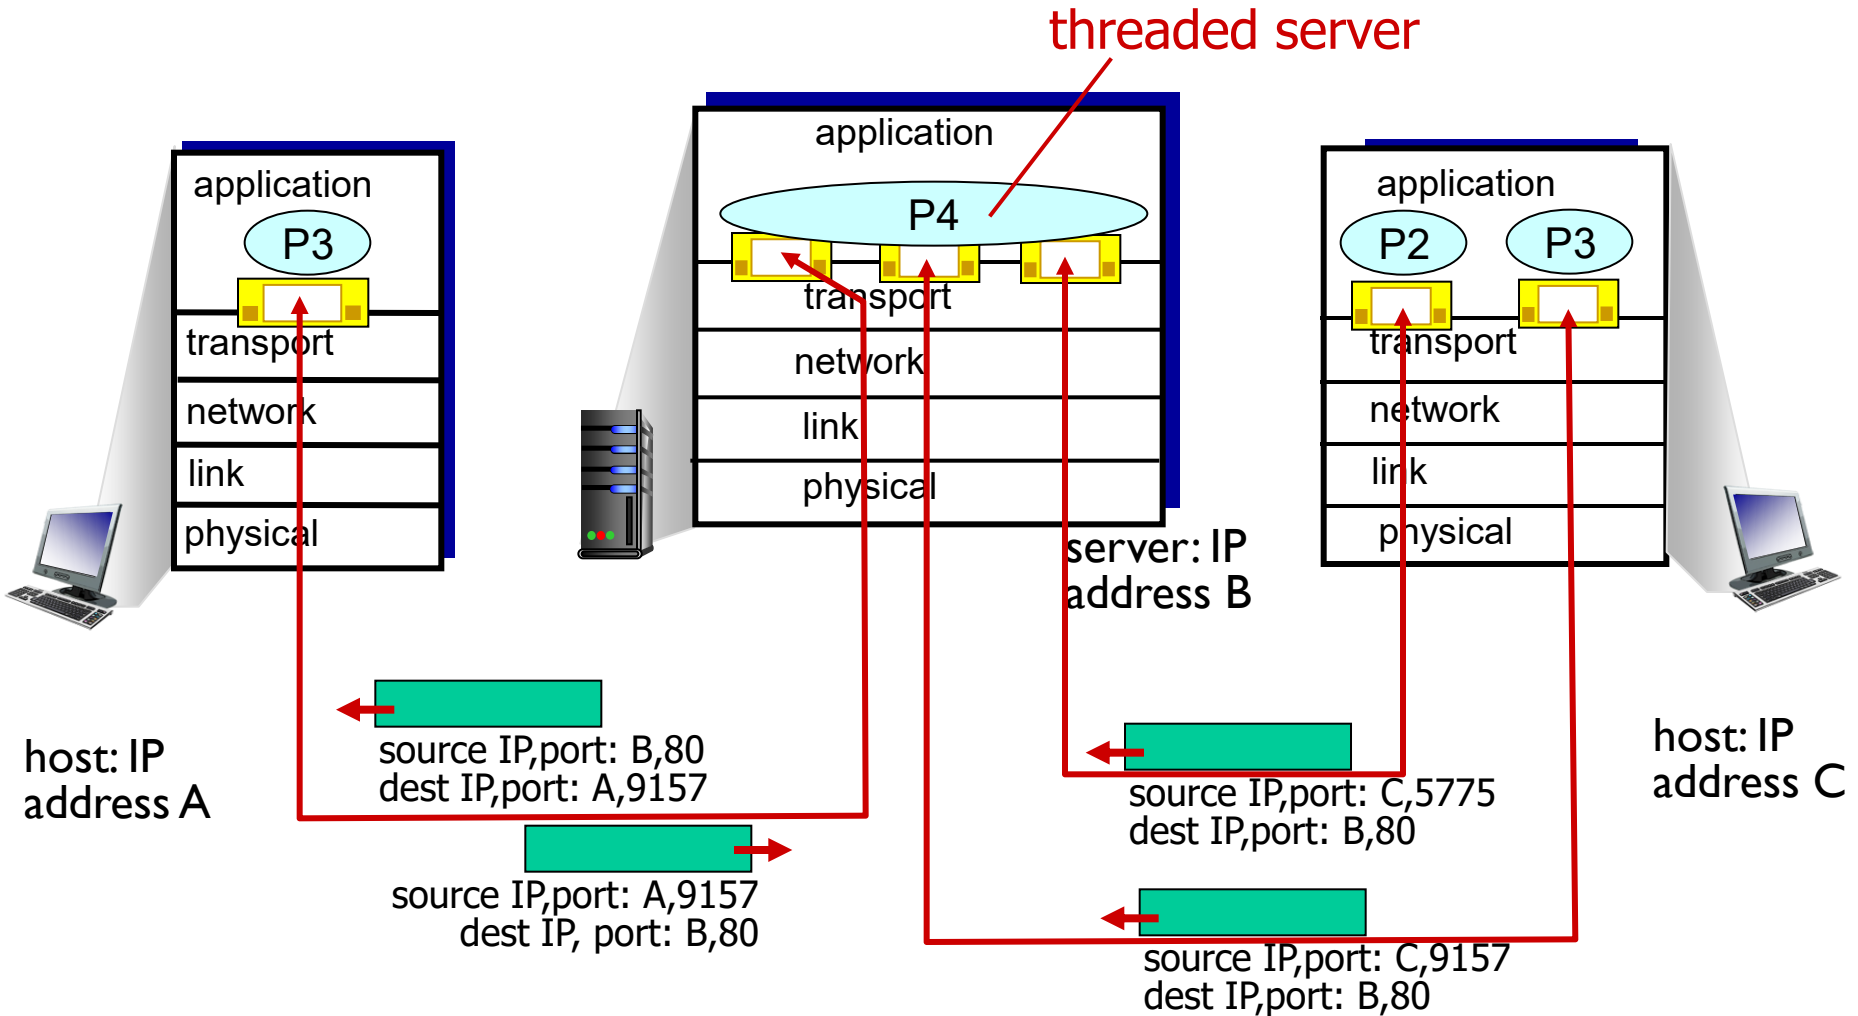
- ❖ TCP socket identified by a 4-tuple:
  - source IP address
  - source port number
  - dest IP address
  - dest port number
- ❖ De-mux: receiver uses all four values to direct segment to appropriate socket
- ❖ server host may support many simultaneous TCP sockets:
  - each socket identified by its own 4-tuple
- ❖ web servers have different sockets for each connecting client
  - non-persistent HTTP will have different socket for each request

# Connection-oriented De-mux – Example



three segments, all destined to IP address: B,  
dest port: 80 are demultiplexed to *different* sockets

# Connection-oriented De-mux – Example



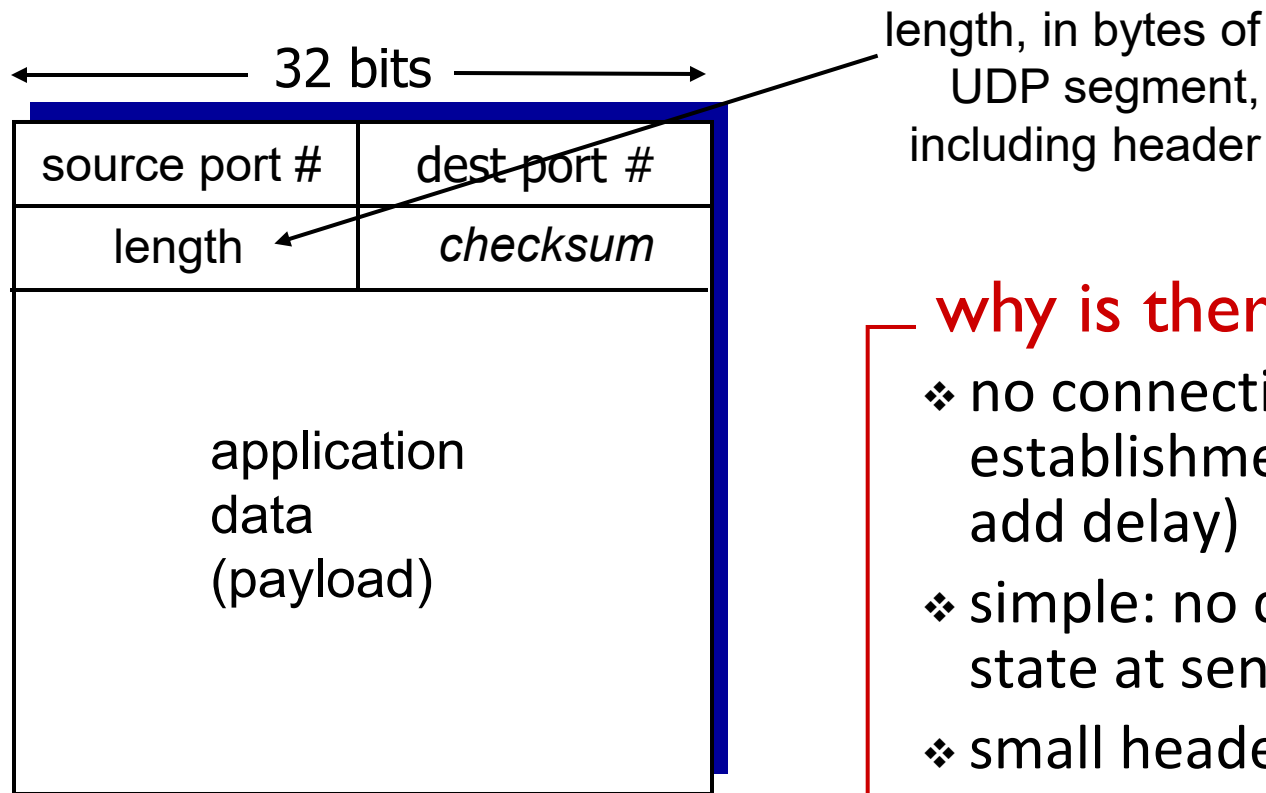
# UDP: User Datagram Protocol [RFC 768]

- “*no frills*,” “*bare bones*”  
Internet transport protocol
  - “*best effort*” service, UDP segments may be:
    - lost
    - delivered out-of-order to app
  - *connectionless*:
    - no handshaking between UDP sender, receiver
    - each UDP segment handled independently of others
- ❖ UDP use:
    - streaming multimedia apps (loss tolerant, rate sensitive)
    - DNS
    - SNMP
  - ❖ reliable transfer over UDP:
    - add reliability at application layer
    - application-specific error recovery!

*What are two main functions does UDP add to IP?*



# UDP: Segment Header



UDP segment format

## why is there a UDP?

- ❖ no connection establishment (which can add delay)
- ❖ simple: no connection state at sender, receiver
- ❖ small header size
- ❖ no congestion control: UDP can blast away as fast as desired

# UDP Checksum

*Goal:* detect “errors” (e.g., flipped bits) in transmitted segment

## sender:

- treat segment contents, including header fields, as sequence of 16-bit integers
- checksum: addition (1-complement sum) of segment contents
- sender puts checksum value into UDP checksum field
  - can be all 0s: no checksum

## receiver:

(If the received checksum is not all 0s)

- compute checksum of received segment
- check if computed checksum equals checksum field value:
  - NO - error detected
  - YES - no error detected.  
(*But maybe errors nonetheless?* More later)
- ....

*What if the sender wants to add checksum and the calculated checksum is all 0s?*

# Internet Checksum Algorithm

1. Break the data (e.g. the UDP message, including header) as a series of 16-bit integers
2. Add them together
3. If there is a carry-over, add 1 to the result
4. Take 1-complement of the result (i.e. reverting all bits) to get the checksum for the data

# Internet Checksum – Example

Message is two 16-bit integers

1	1	1	0	0	1	1	0	0	1	1	0	0	1	1	0
1	1	0	1	0	1	0	1	0	1	0	1	0	1	0	1

Carry-over	1	1	0	1	1	1	0	1	1	1	0	1	1	1	0	1	1
------------	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

1-compliment sum	1	0	1	1	1	0	1	1	1	0	1	1	1	1	0	0
checksum	0	1	0	0	0	1	0	0	0	1	0	0	0	0	1	1